MAE per Parameter

MAE (Mean Absolute Error) measures the average absolute difference between the predicted and actual values. Lower values indicate better performance. A small spread across parameters suggests consistent predictions.
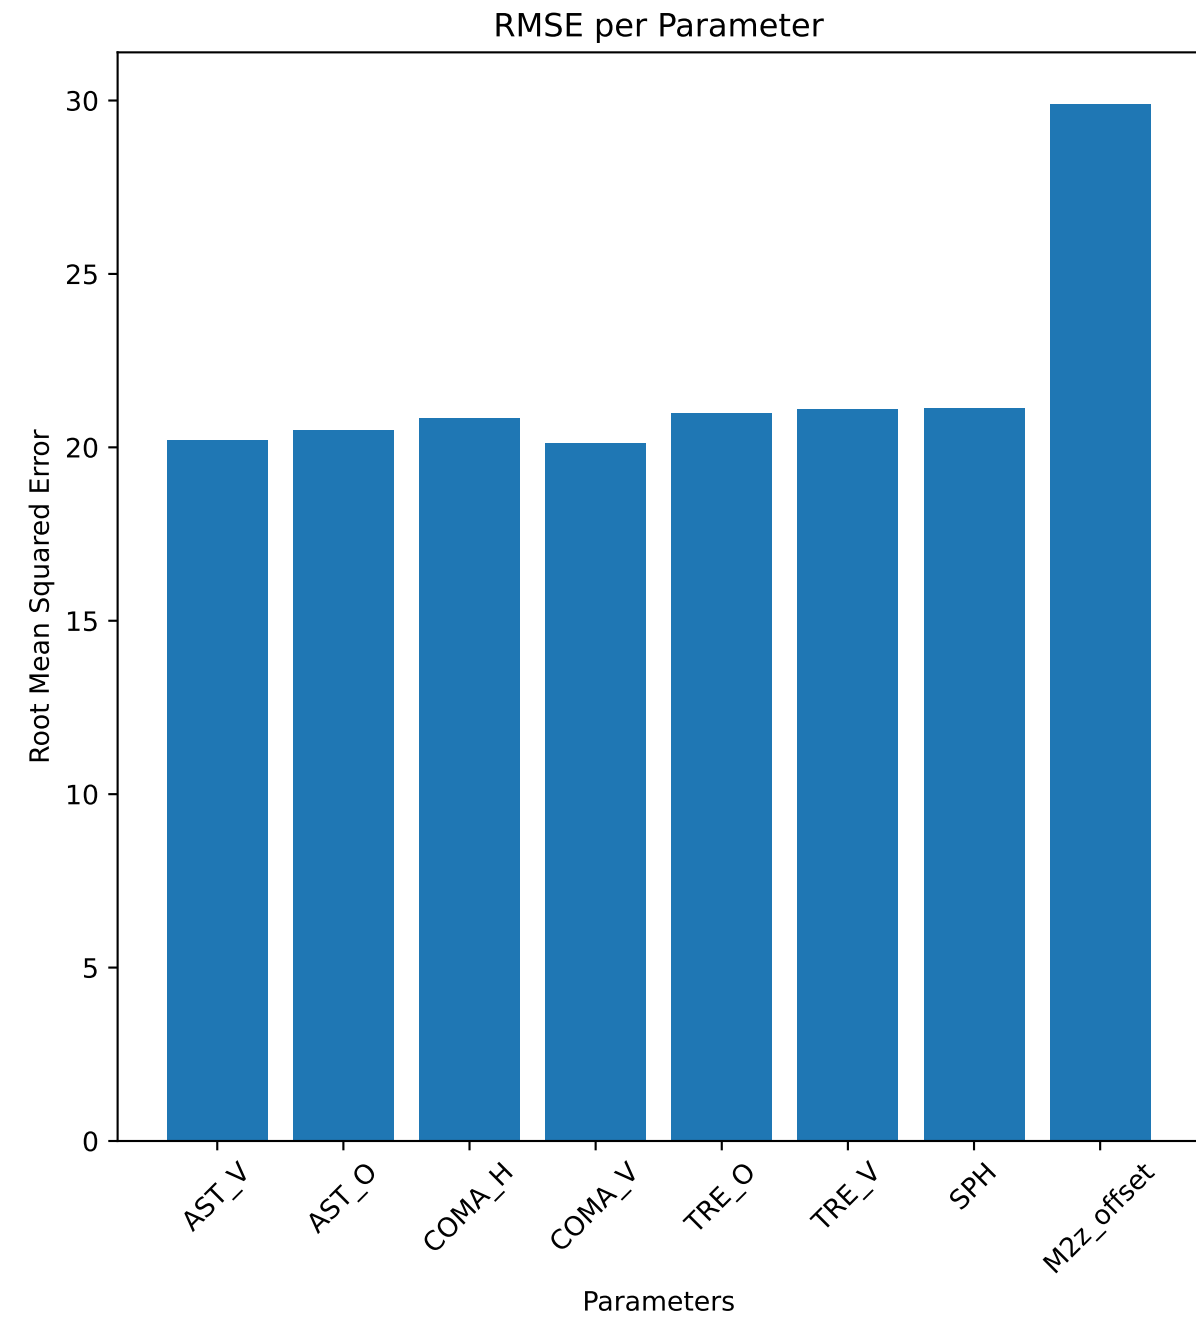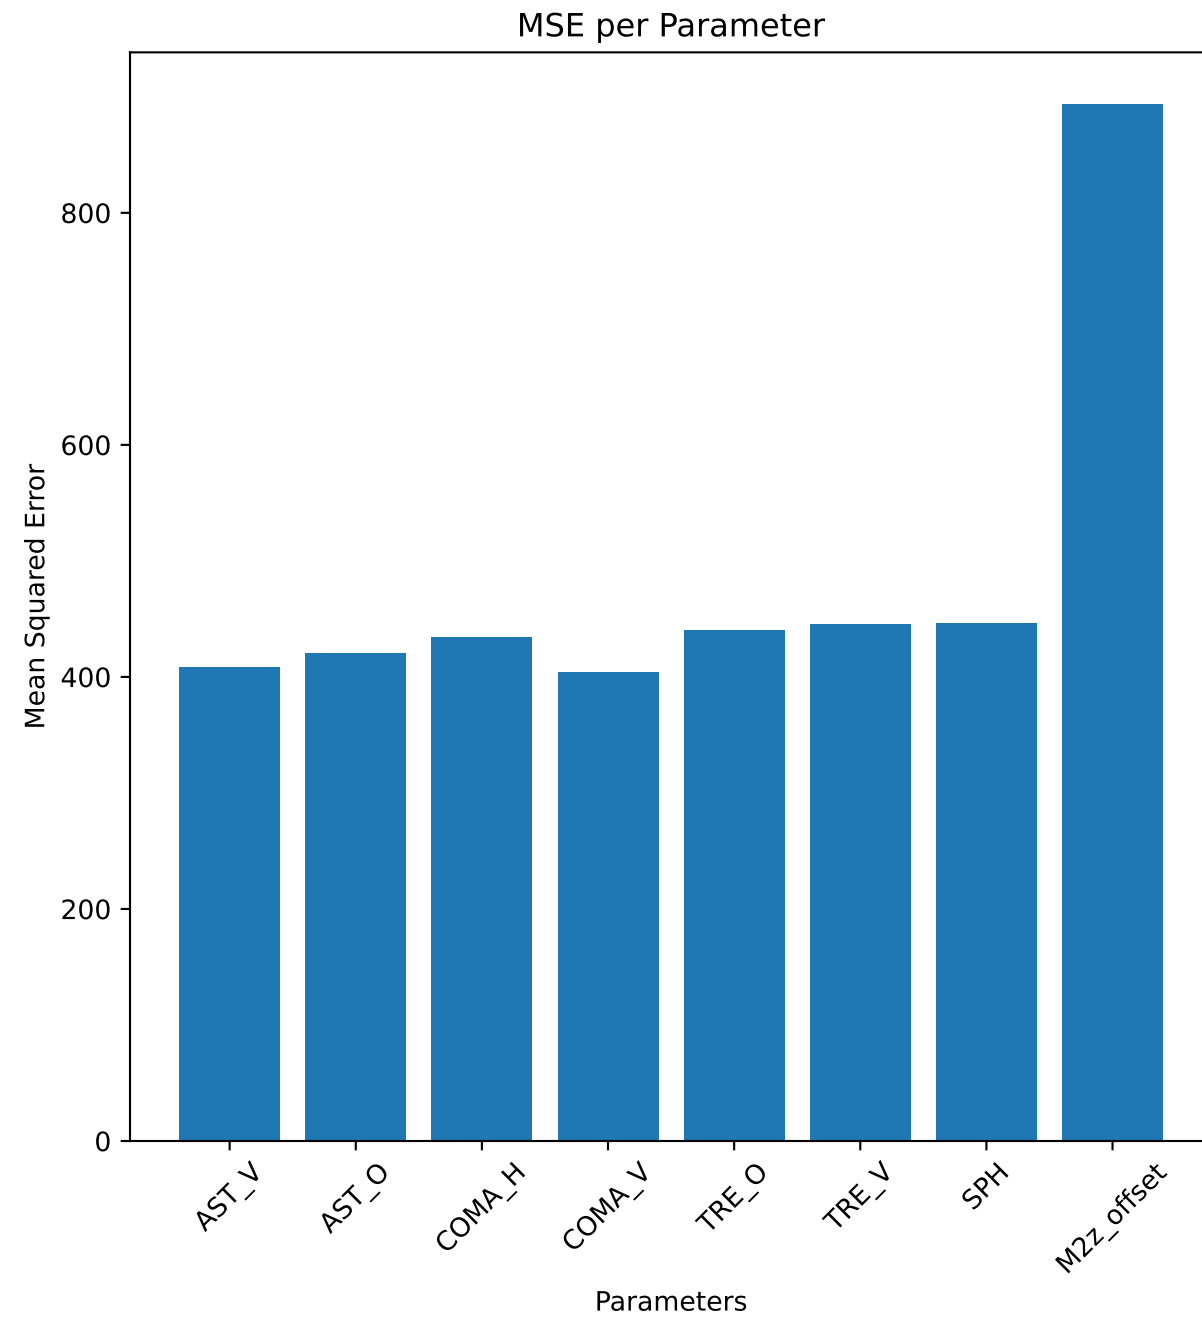
MSE (Mean Squared Error) squares the errors, penalizing larger deviations more than MAE, while RMSE (Root Mean Squared Error) converts this back to the original scale. Lower values indicate a better fit.
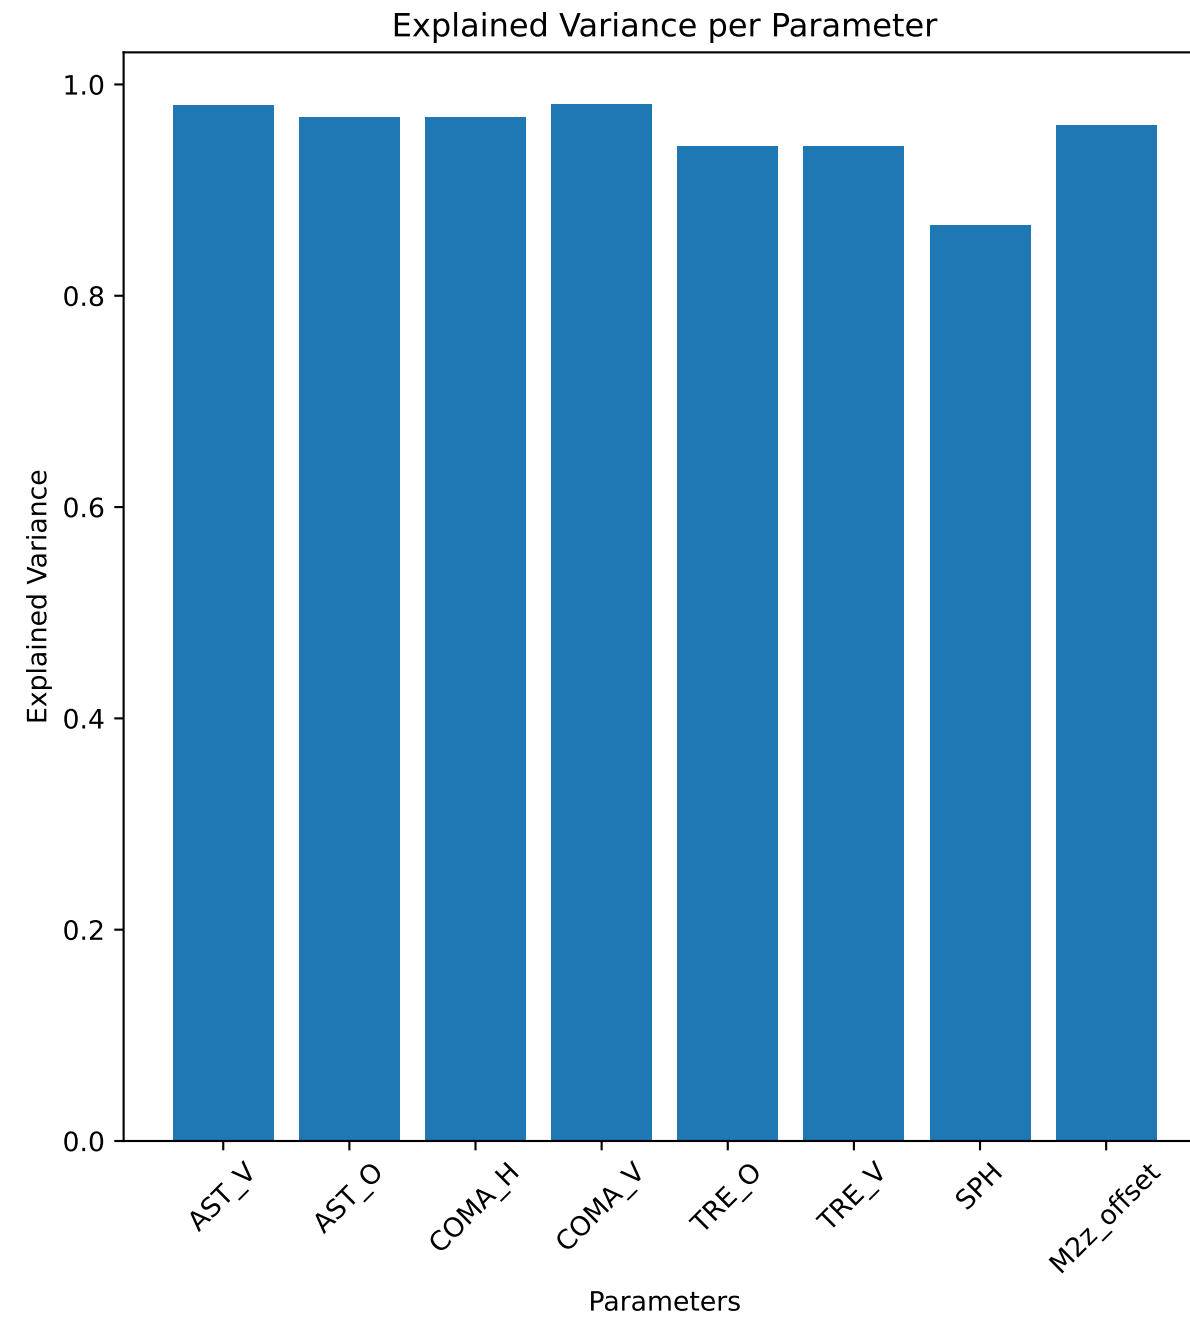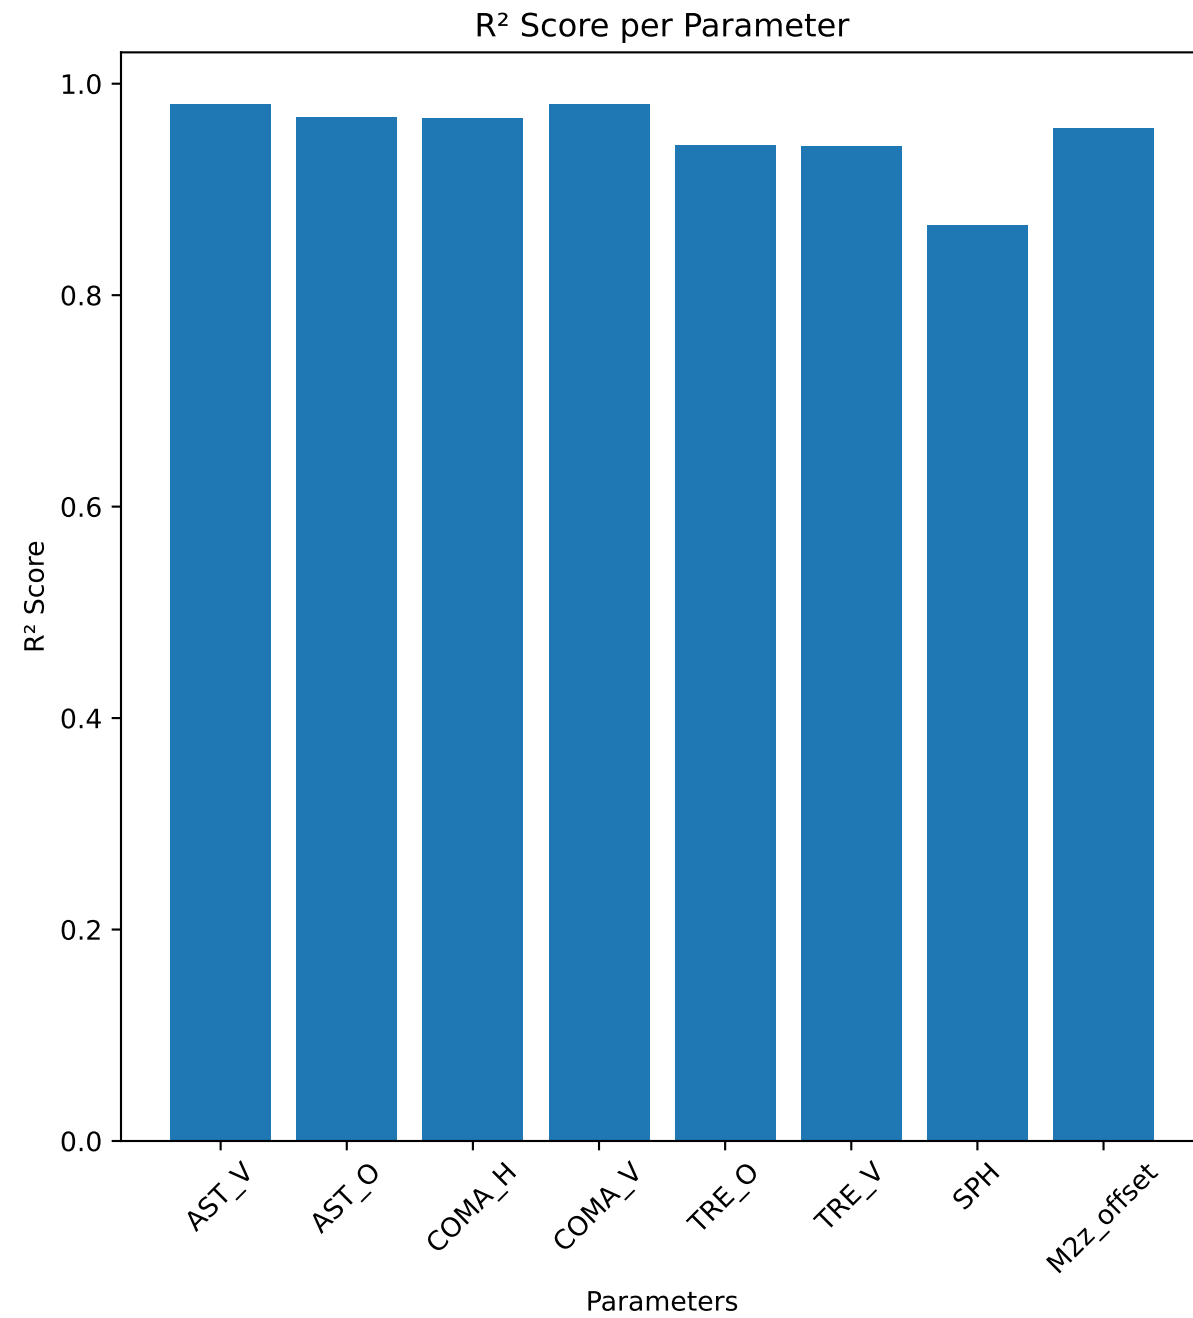
The R² Score indicates the proportion of variance in the data explained by the model, with 1 being ideal. Explained Variance measures how much variability is captured. High values (near 1) are desirable.

## Median Absolute Error per Parameter
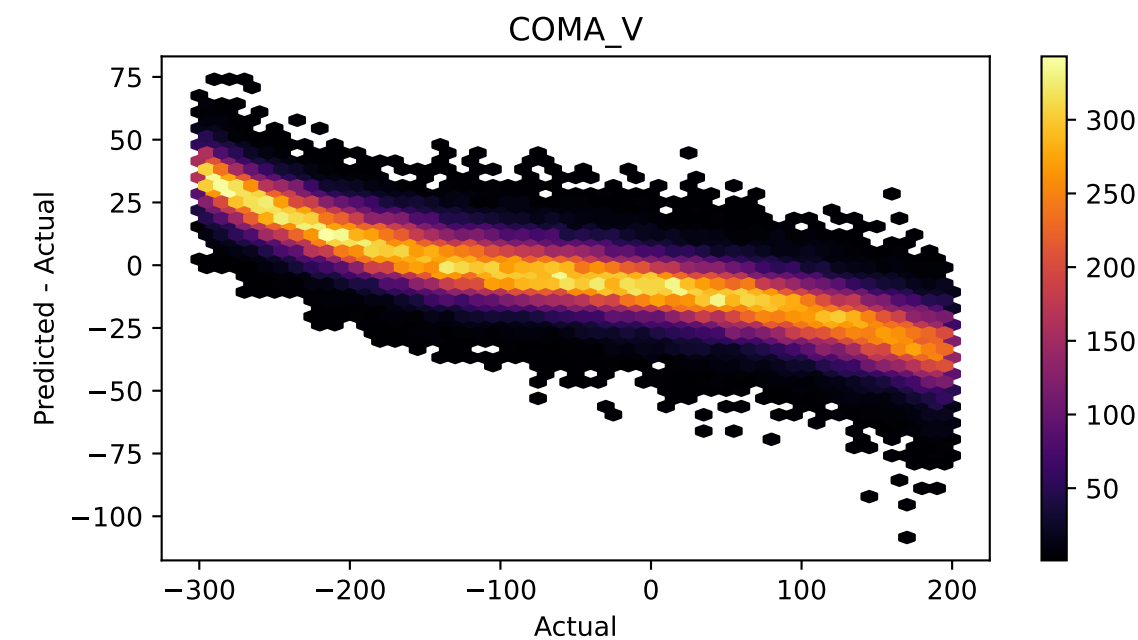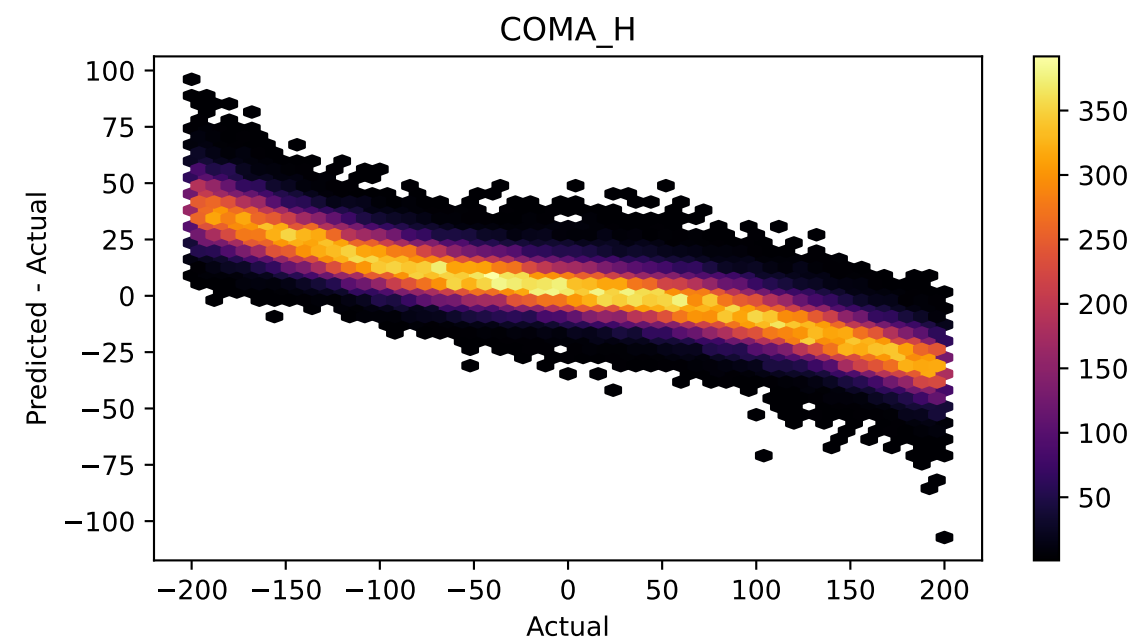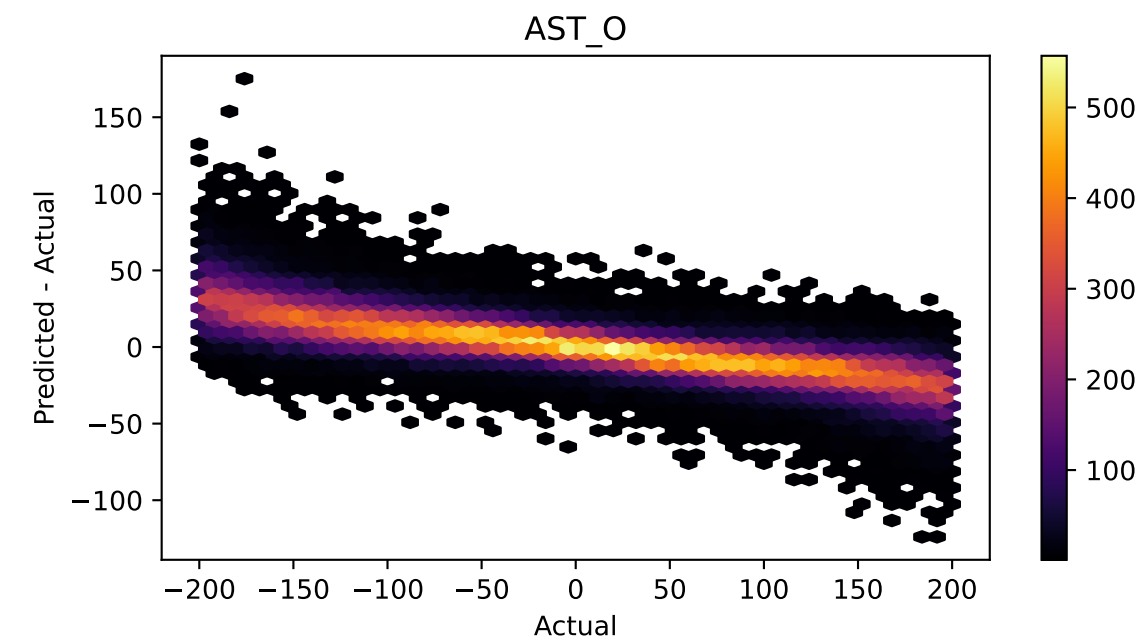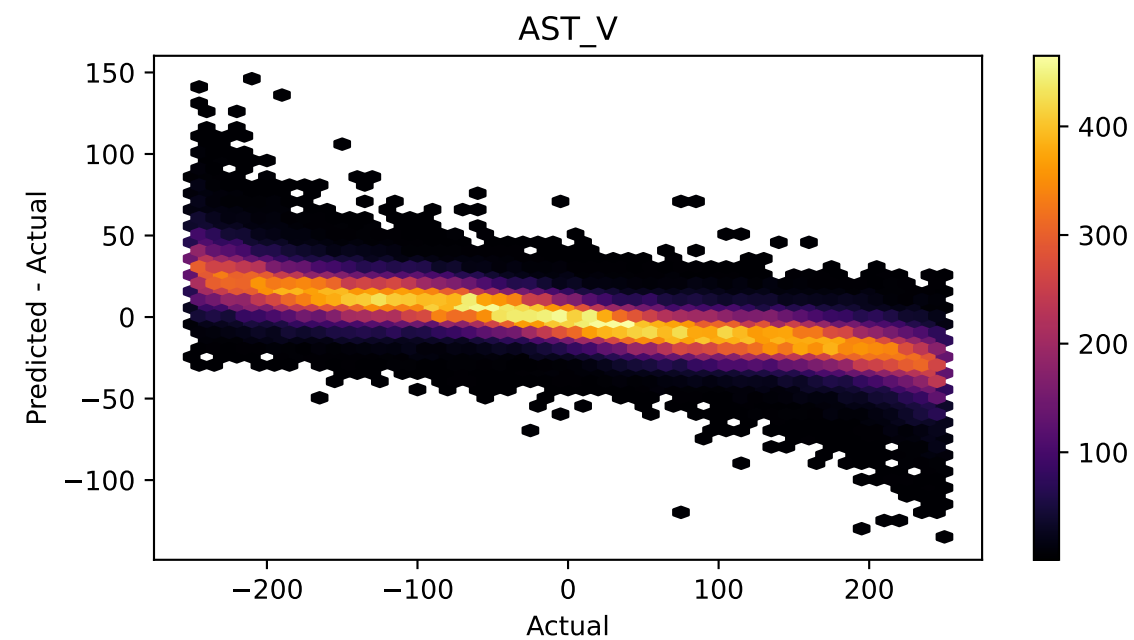
## Maximum Error per Parameter
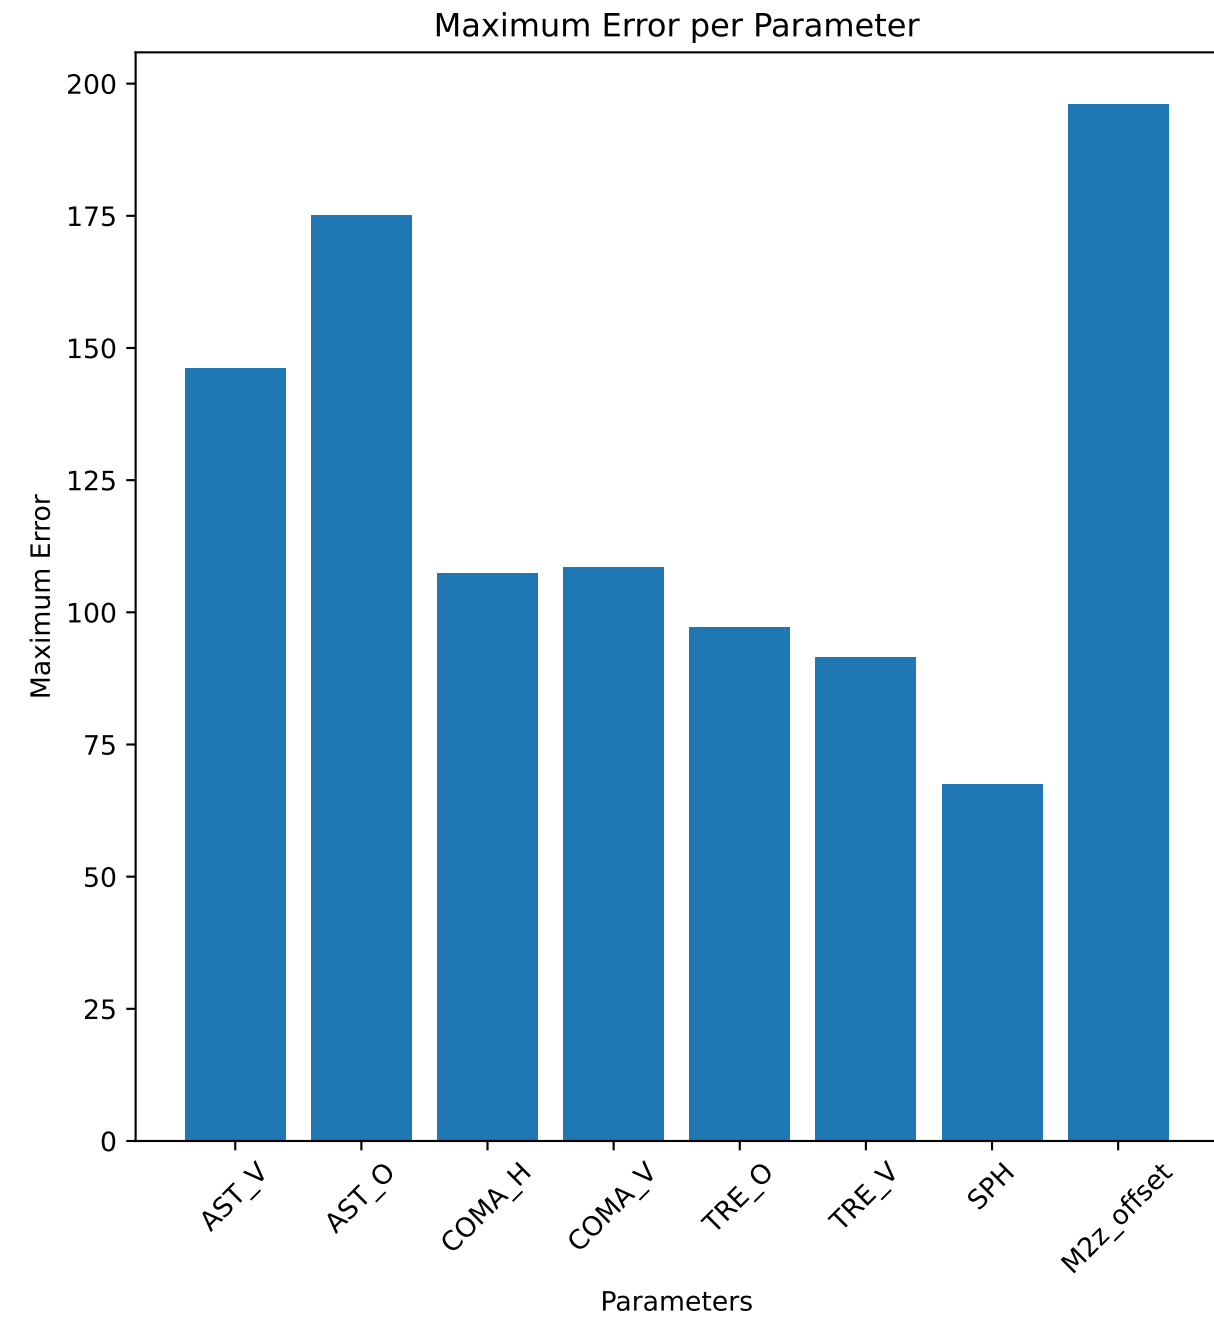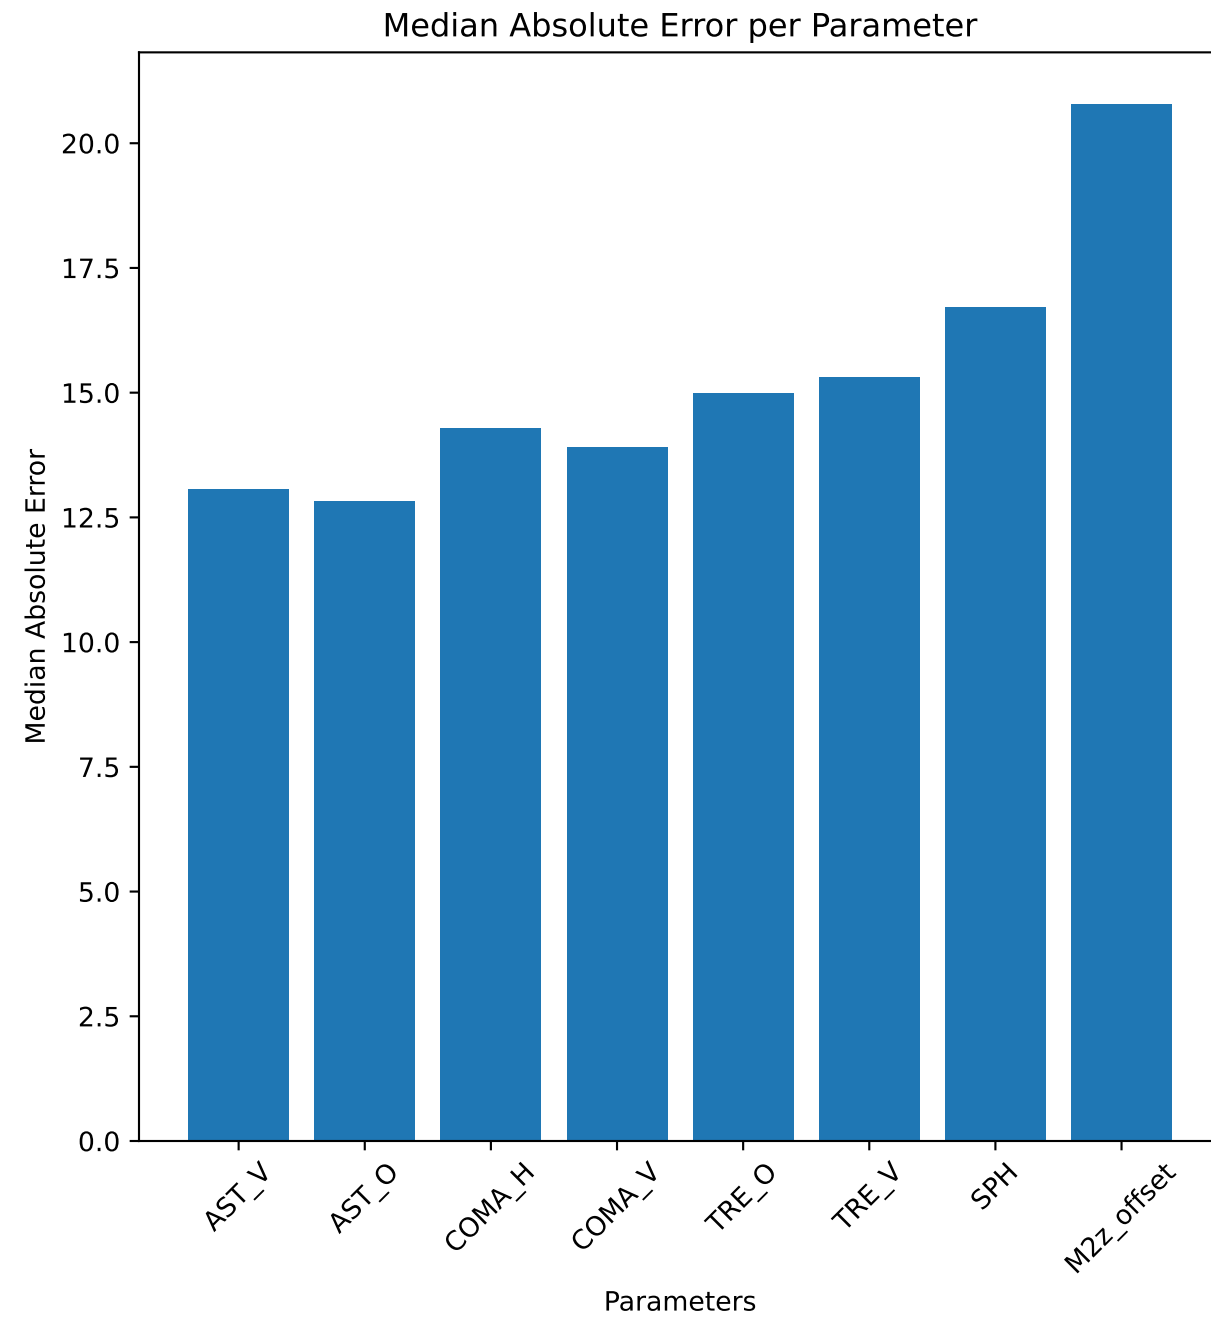
Median Absolute Error provides a robust central tendency of errors, while Maximum Error indicates the worst-case scenario. Ideally, both should be low.
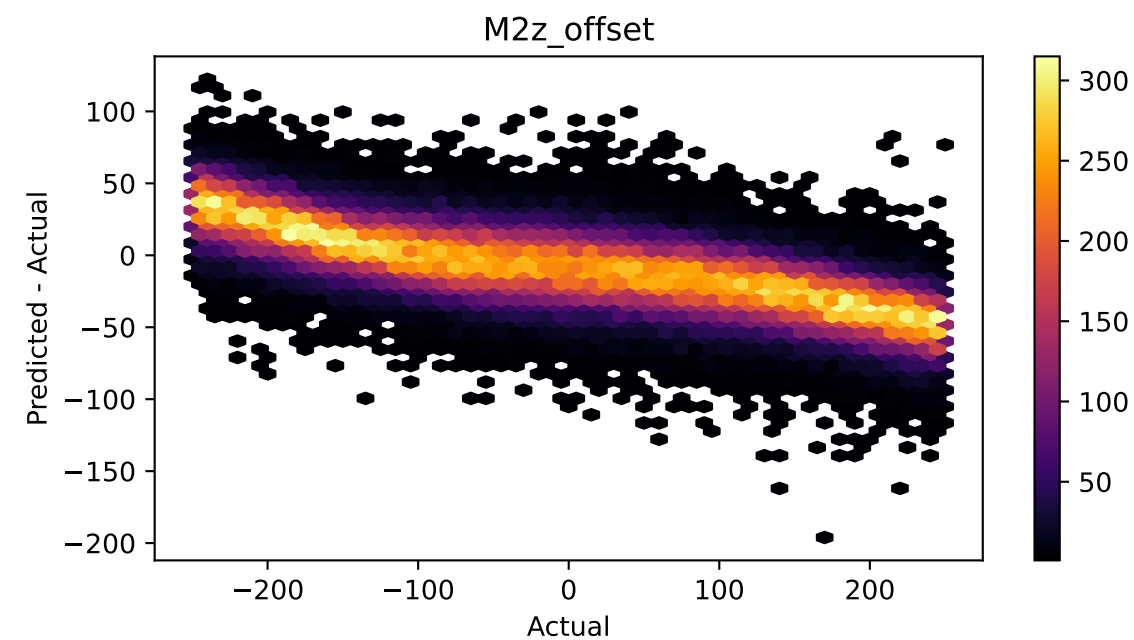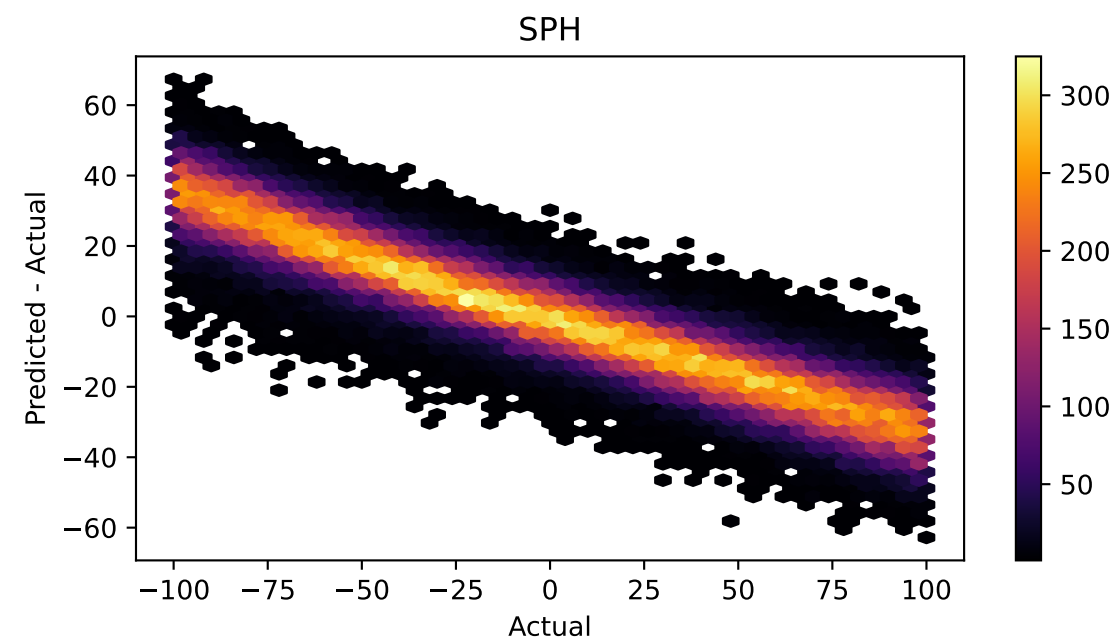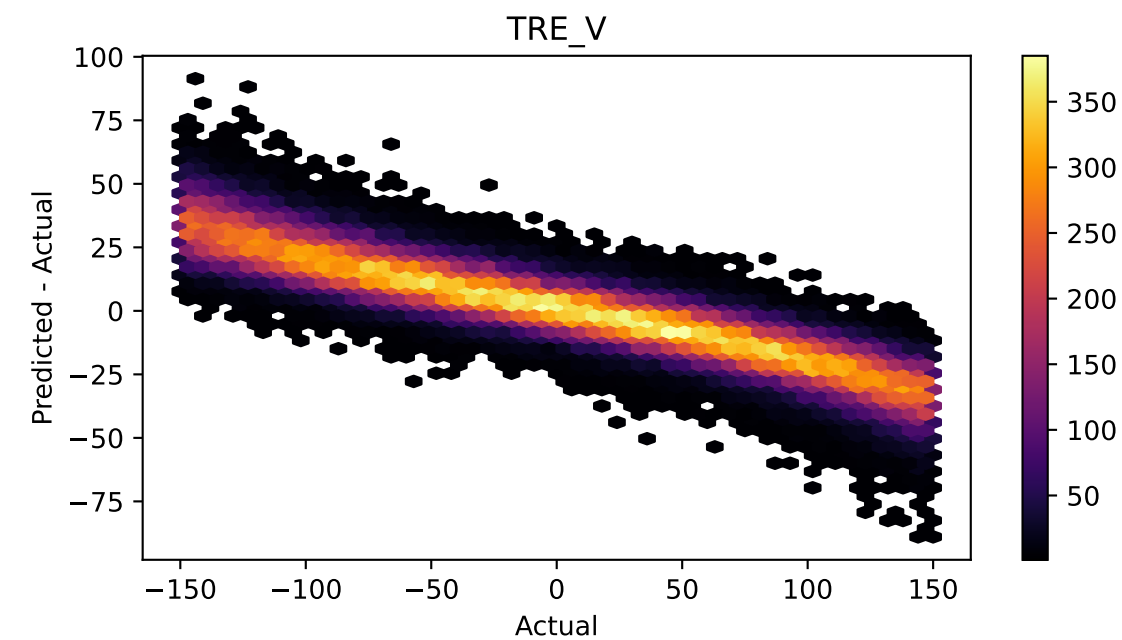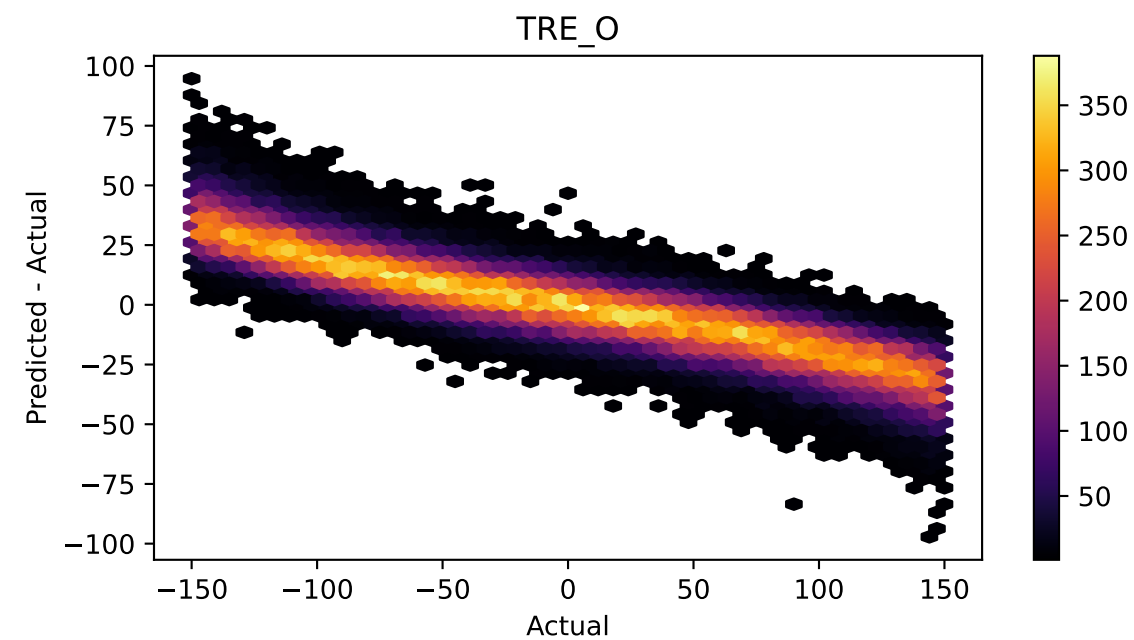
## Distribution of Residual Errors
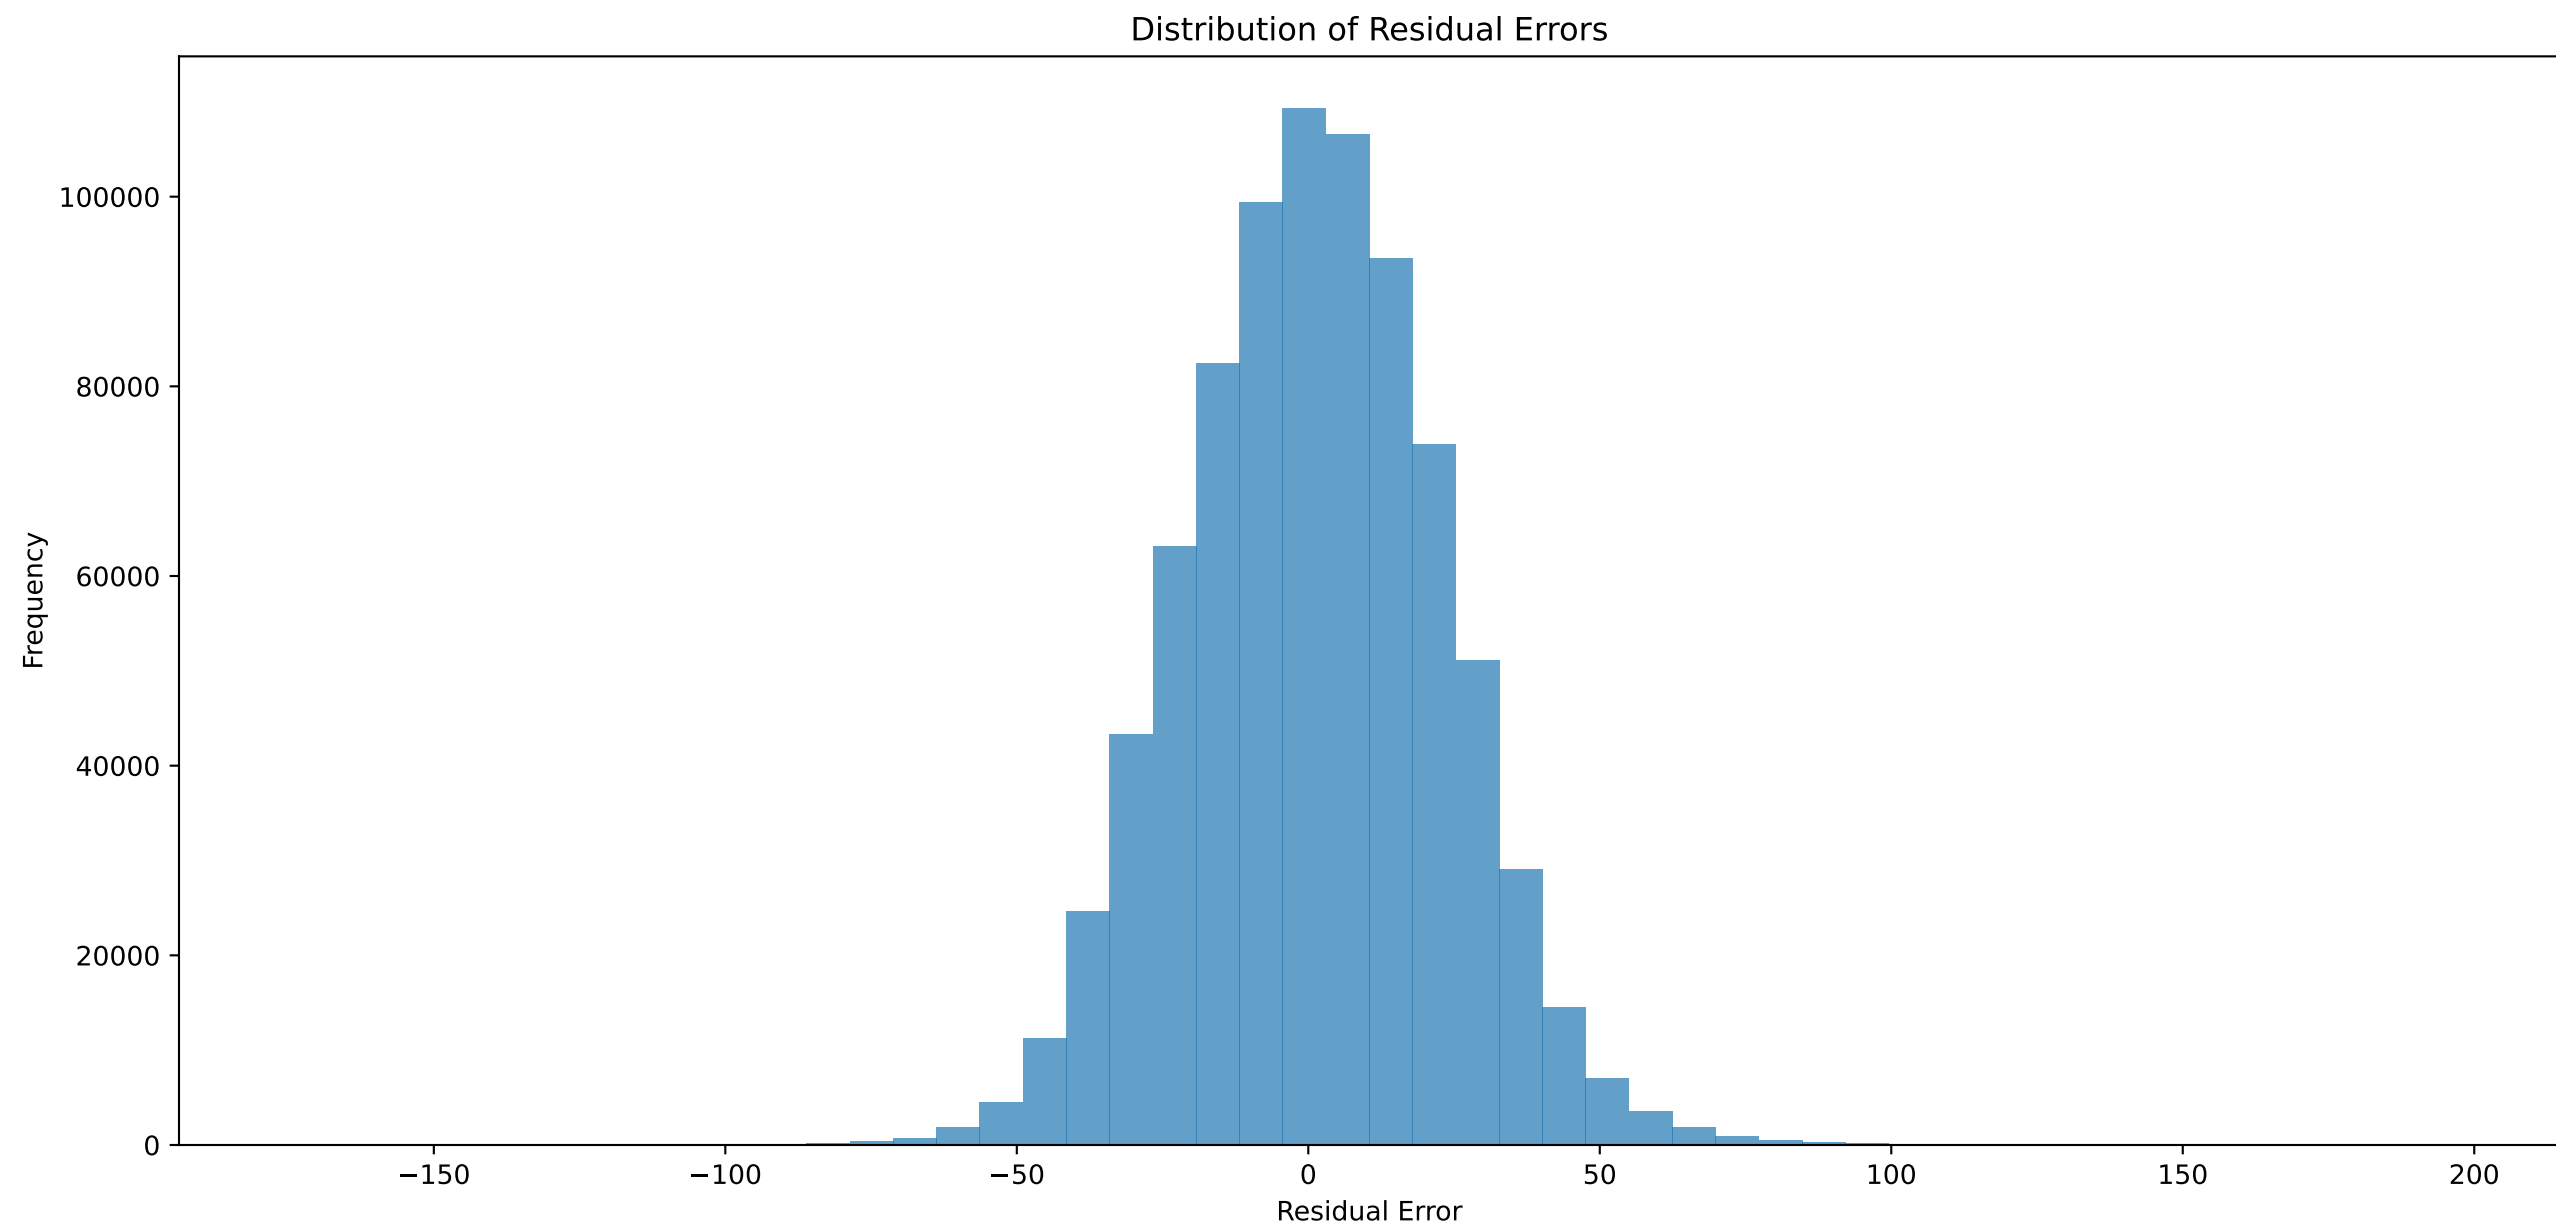
Shows the distribution of (y_true - y_pred). A symmetric distribution centered near zero suggests minimal bias.

Hexbin: Actual vs. (Predicted-Actual)

Each plot compares the actual and residual for a parameter using a hexbin plot. Concentration near zero indicates good predictions.

Hexbin: Actual vs. (Predicted-Actual)

Each plot compares the actual and residual for a parameter using a hexbin plot. Concentration near zero indicates good predictions.

Correlation Heatmap of Predicted Parameters
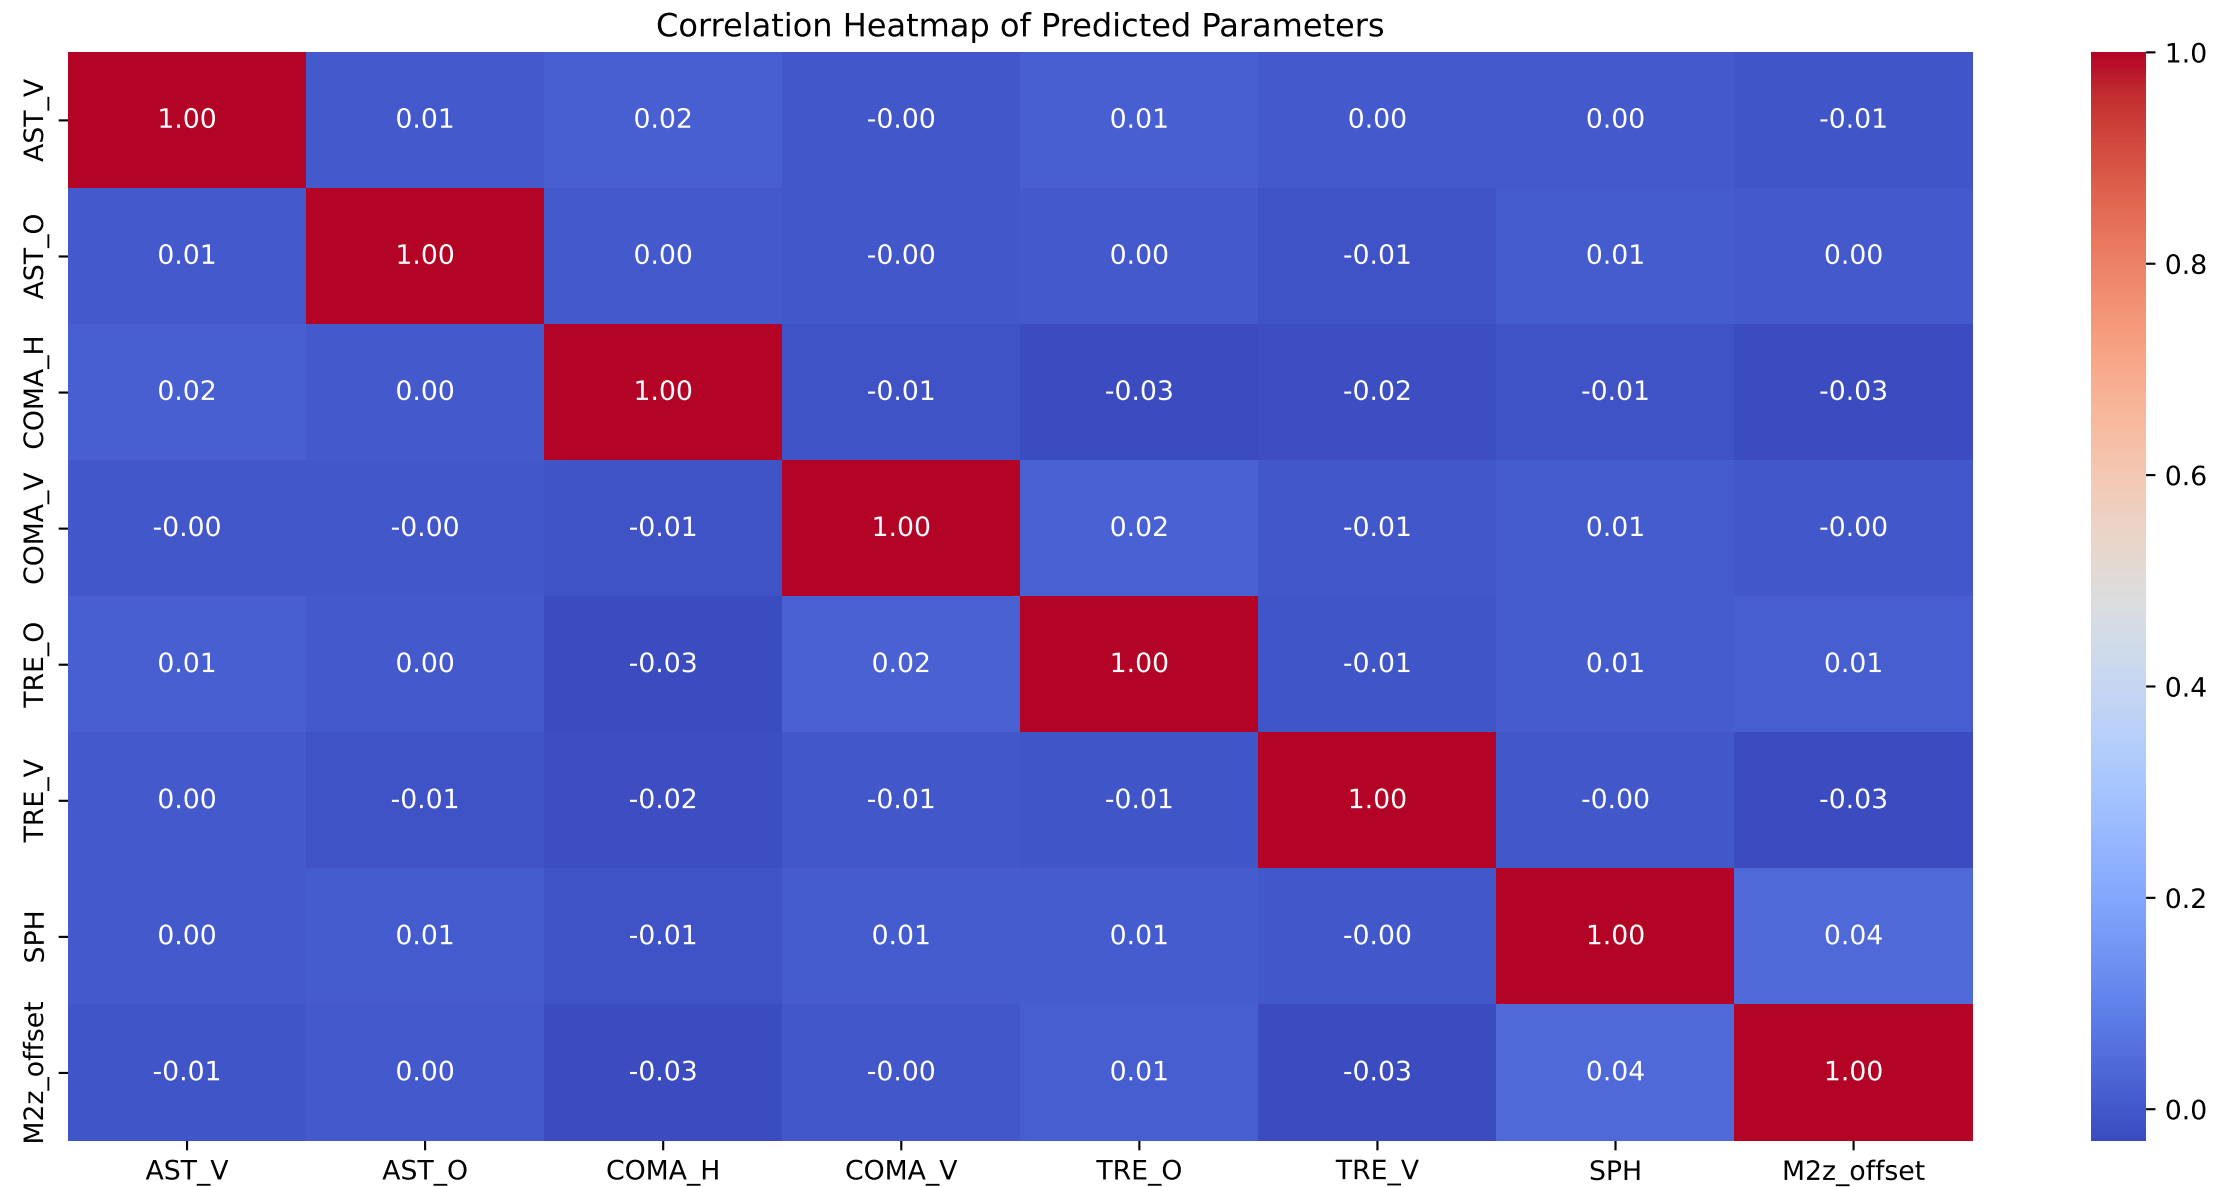
Shows pairwise correlations between predicted parameters. High correlations may reveal underlying couplings or redundancies.